

# Supervised Learning & CNNs

Narada Warakagoda

Outline Architecture

Losses

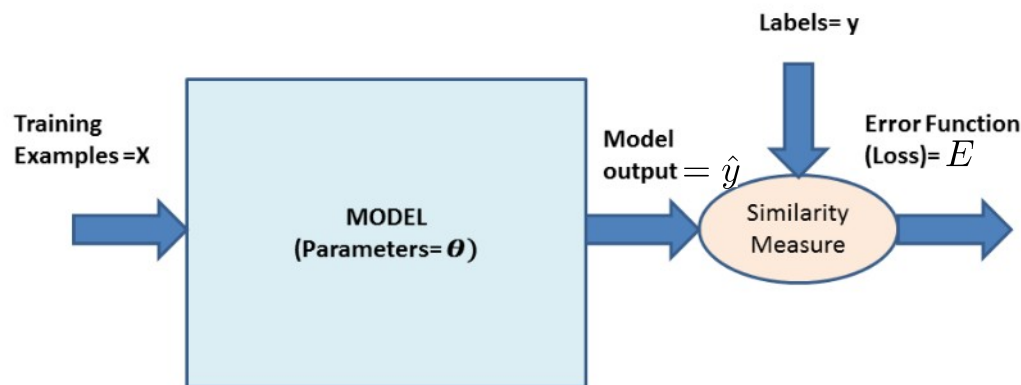
Training

Regularization

1 / 30

# Supervised Learning outline


- Given: A set of training examples (input, label) pairs  $\{ (x,y) \}$
- Find: Model parameters  $\theta$  such that the similarity between the model output  $\hat{y}$  and the labels  $y$  are maximized. Similarity is expressed as a loss/error/cost function.



# Main aspects of supervised deep learning


- What is the model architecture?
- What is the loss function?
- How do we update the model parameters?
- How do we maintain the generalization ability of the model?

# Neuron



Biological Neuron


$$y = f(\sum_i w_i x_i + b) = f(\mathbf{x} \cdot \mathbf{w}^T + b)$$




Mathematical Model

Illustrations from <http://cs231n.github.io/neural-networks-1/>


# Activation Functions



Sigmoid




Tanh Sigmoid



Rectified Linear (ReLU)

# Fully Connected (Dense) Network



- Each layer performs  $y = f(\mathbf{x}\mathbf{W} + \mathbf{b})$
- Single hidden layer can approximate any function (Universal approximation theorem)
- Number of parameters can grow quickly

# Convolutional Neural Network (CNN)

- Can be seen as:
  - Crude model of human visual cortex
  - Generalization of Gabor filters
  - Generalization of template matching
  - Way of parameter sharing and hence parameter reduction

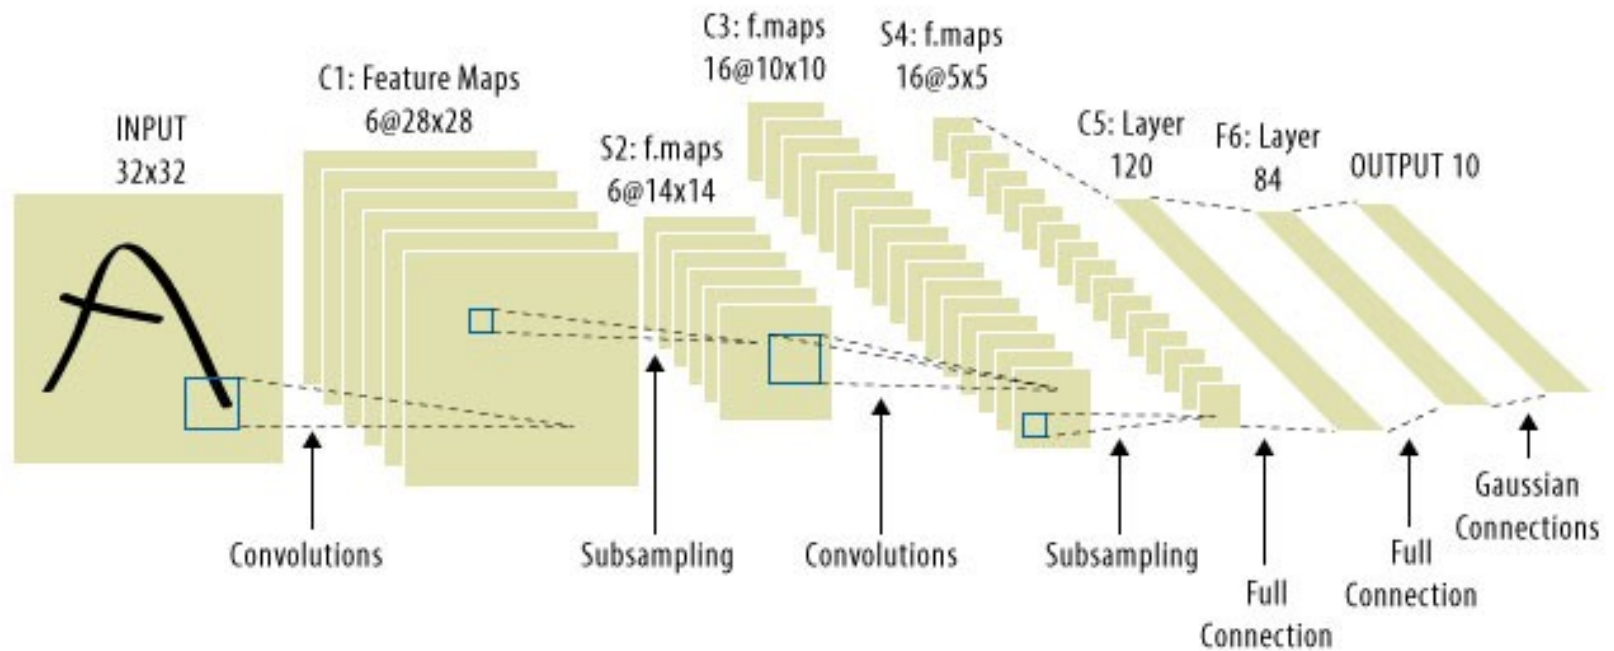
# CNN as glorified template matching

- Try to match template as each location by sliding it over the input image



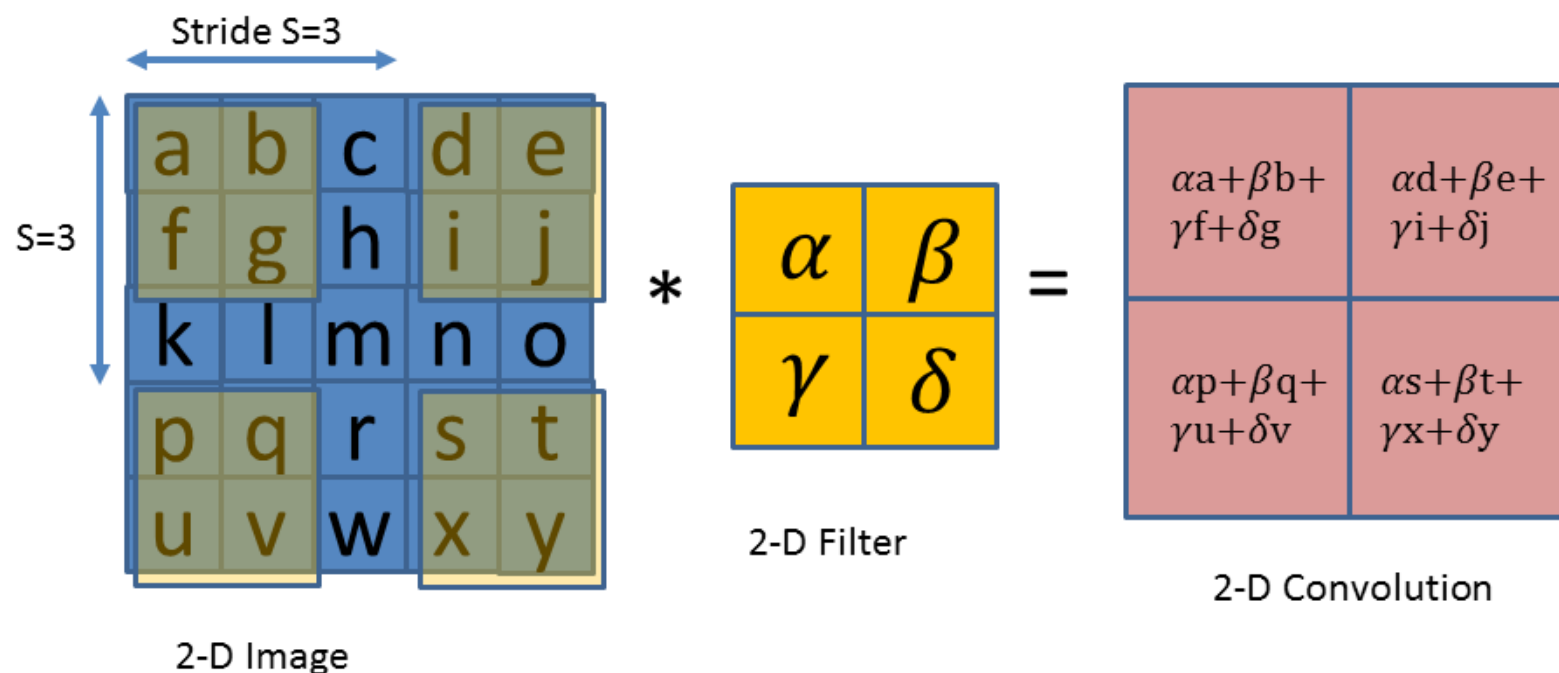


# A typical CNN



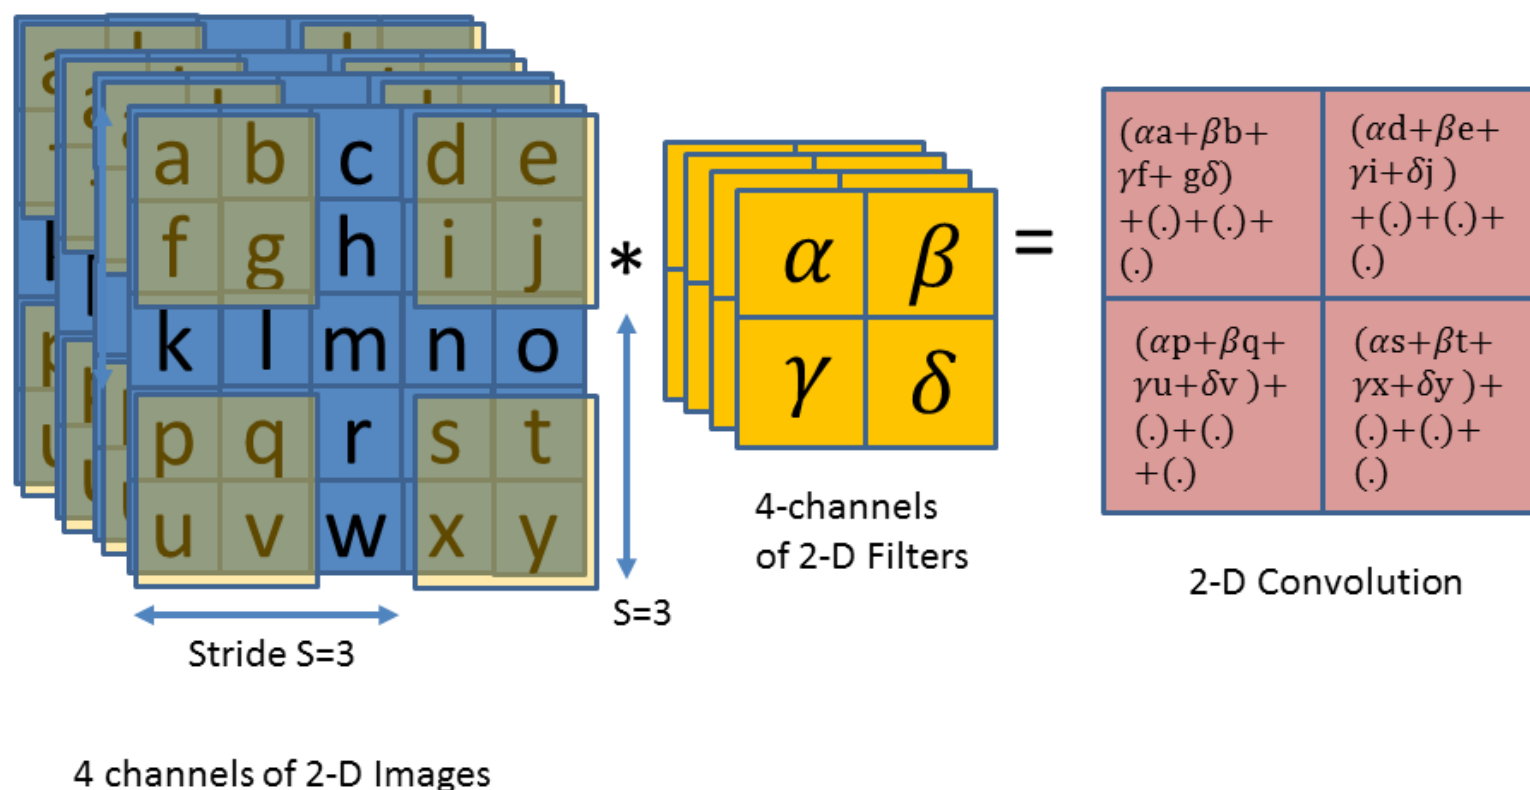
# Convolutional Layer

- Convolution of 2D-inputs (eg: single channel images)



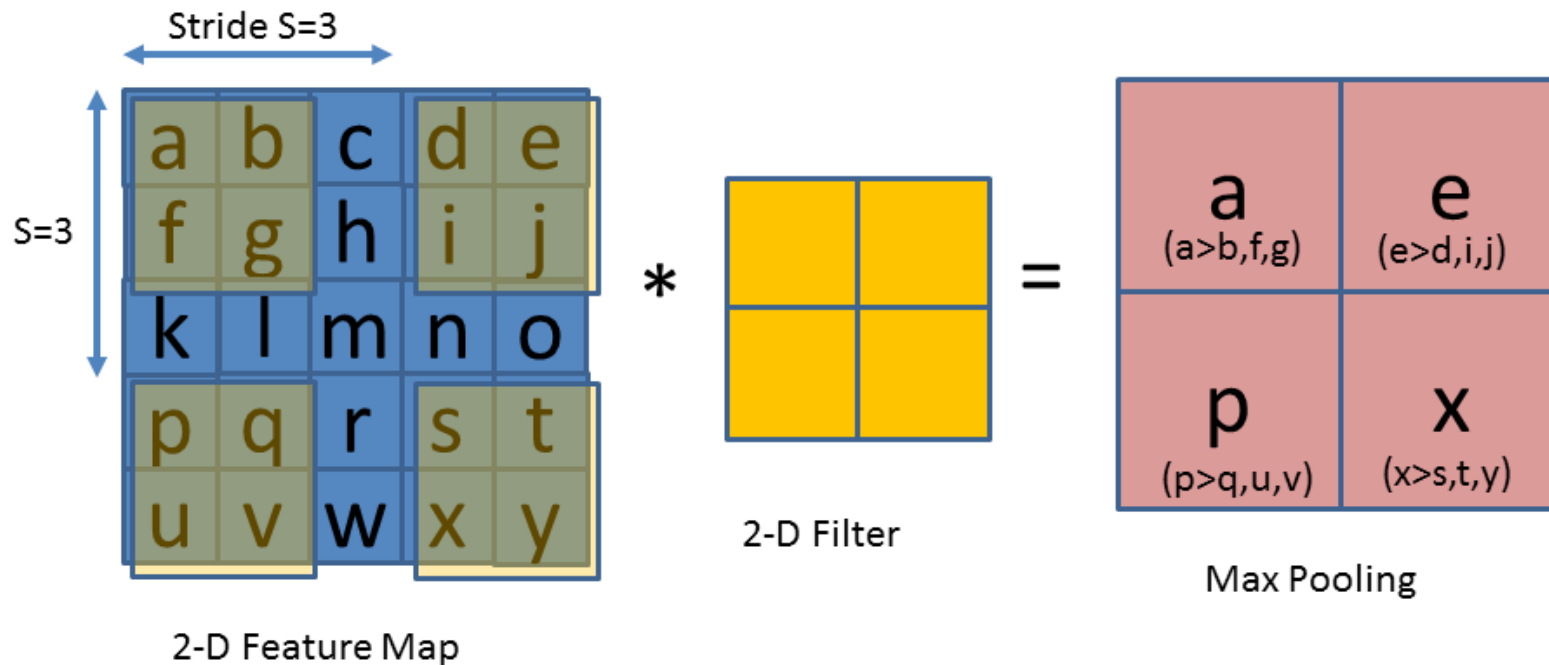
# Convolutional Layer

- Convolution with 3D inputs (eg: multiple channel images or feature maps)



# Sub-sampling layer

- Max-pooling operation




# Padding

- Convolution and sub-sampling (max-pooling) lead to an output dimension  $N_o = \frac{N_i - F}{s} + 1$  where  $N_o$  = input dimension,  $F$  = Filter dimension and  $s$  = stride
- Output dimension can decrease even when stride  $s = 1$ .
- Apply zero padding around the image to prevent such reduction of dimensions.


# Padding

- Tensorflow padding options illustrated using 1D signals

VALID Padding



SAME Padding




# Loss functions

- Also called Cost/Error/Objective functions
- A loss function measures the similarity between model output  $\hat{y} = f_{\theta}(\mathbf{X})$  and the corresponding label  $y$
- Common loss functions:
  - Mean Squared Error (MSE)
  - (Categorical) Cross Entropy (CE)

# Mean Squared Error (MSE)

- $E_{MSE} = \frac{1}{N} \sum_i (y_i - \hat{y}_i)^2$
- Suitable for regression (i.e. predict  $y_i$ )
- It can be shown that MSE is equivalent to the conditional maximum likelihood of labels, when the prediction error  $\epsilon_i$  has a zero mean Gaussian distribution.

$$\arg \min_{\theta} \{E_{MSE}\} = \arg \min_{\theta} \{-\sum_i \log p(y_i | \mathbf{X}_i)\}$$





# Cross Entropy (CE)

- Suitable for classification (i.e. predicting the class probabilities of  $K$  given classes)

- Labels can be:

- One-hot encoded or class probability label:

$$[y_i(1), y_i(2), \dots, y_i(K)]$$

- Sparse: index of the only correct class  $j$  is given, i.e.  $y_i = j$

# Cross Entropy (Ctd)


- Model outputs a categorical distribution i.e. probability of each class  $p(c(k)|\mathbf{X}_i)$ , where  $i$  is the sample index and  $c(k)$  is the class of index  $k$
- Cross Entropy:

- One-hot encoded or class probability labels:

$$E_{CE} = -\frac{1}{N} \sum_i \sum_k y_i(k) \log p(c(k)|\mathbf{X}_i)$$

- Sparse labels: Assume the correct class is  $y_i$

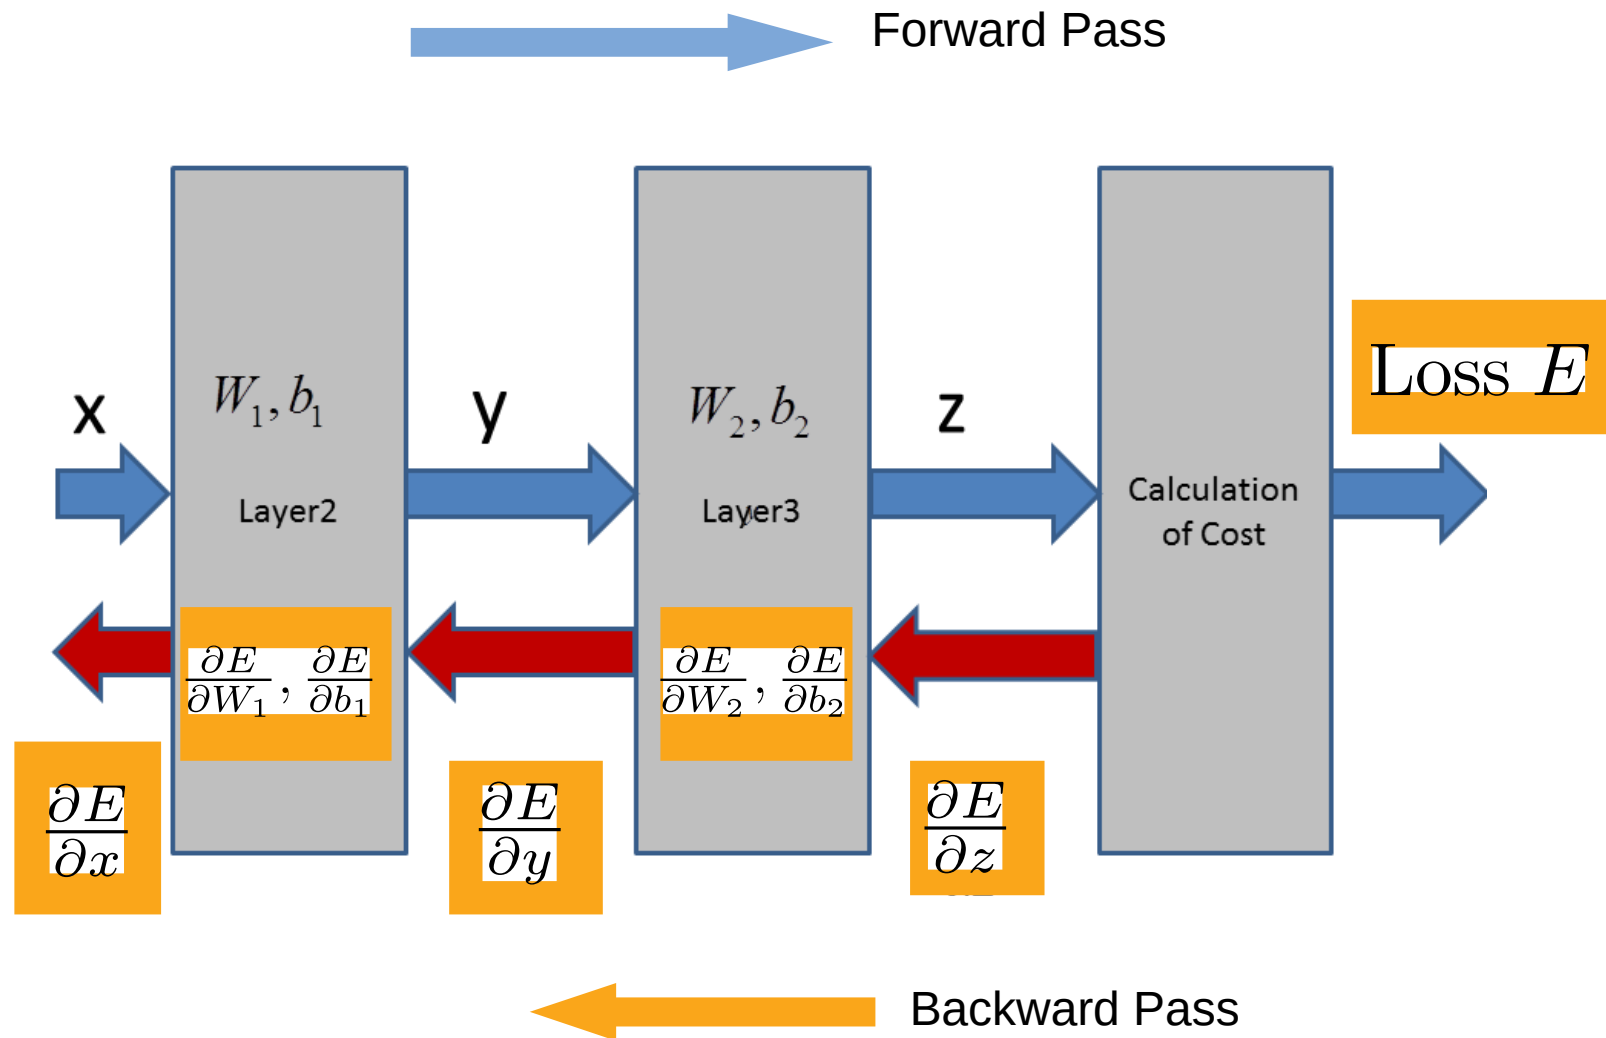
$$E_{CE} = -\frac{1}{N} \sum_i \log p(c(y_i)|\mathbf{X}_i)$$



# Training

- The process of fitting the model parameters to the given training data
  - This is the same as optimizing the loss function with respect to model parameters
- Parameters are updated using gradient descent
  - Plain gradient descent:  $\theta_{t+1} = \theta_t - \epsilon \frac{\partial E}{\partial \theta_t}$
  - More fancy algorithms (ADAM, RmsPROP, and many more)
- Estimation of gradients  $\frac{\partial E}{\partial \theta_t}$  is the core problem of training
  - Back-propagation algorithm

# Back-Propagation







# Model Parameter Update

- Batch gradient descent:
  - Gradients are accumulated for the whole training set, and do the update
- Stochastic gradient descent:
  - No gradient accumulation, update after processing each (randomly selected) sample
- Mini-batch gradient descent:
  - Divide the training set into several mini-batches, accumulate gradient and update after each (randomly selected) mini-batch

# Regularization

- Network should perform well on unseen data
  - Generalization ability
- Too much adaptation to training data (over-fitting) harms generalization ability.
- Regularization prevents over-fitting



# Regularization vs Optimization

- Optimization: Try to reduce training loss
  - Often test loss is reduced as a side effect
  - Test loss may not be the minimum if over-fitting occurs
- Regularization: Try to reduce the test loss
  - Regularization imposes “restrictions” on training.
  - Minimum test loss is not necessarily corresponding to the minimum training loss

# Popular regularization techniques in deep learning

- Early stopping in training
- Data shuffling
- Data Augmentation
- Weight decay
- Dropout
- Batch Normalization




# Weight Decay

- Modify the loss function
  - Add a penalty term for high parameter values
  - $E_{reg}(\theta) = E(\theta) + \lambda ||\theta||^2$
  - $\lambda$  is a positive constant and  $||\cdot||$  is Euclidean norm of the parameter vector  $\theta$

# Dropout

- Randomly remove neurons from a given layer every time a sample is processed
- Removal percentage is a parameter of Dropout
- Typically done only in training (i.e. NOT in testing)



# Batch Normalization (BN)

- Normalize the input of a neuron (or inputs to a set of related neurons) with the mean and standard deviation of the batch.
- This will prevent “covariate shift” (i.e. large changes of input distribution from one batch to another)
- Initially proposed as a method for improving the training speed
- BN has a regularization effect as well.

# BN for image data

$$\mu_d \leftarrow \frac{1}{N * H * W} \sum_{i=1}^N \sum_{h=1}^H \sum_{w=1}^W x_{i,h,w,d}$$

$$\sigma_d^2 \leftarrow \frac{1}{N * H * W} \sum_{i=1}^N \sum_{h=1}^H \sum_{w=1}^W (x_{i,h,w,d} - \mu_d)^2$$

$$\hat{x}_{i,h,w,d} \leftarrow \frac{x_{i,h,w,d} - \mu_d}{\sqrt{(\sigma_d^2 + \epsilon)}}$$

$$y_{i,h,w,d} \leftarrow \gamma \hat{x}_{i,h,w,d} + \beta$$

- $x_{i,h,w,d}$  is an element of the input tensor of size  $N \times H \times W \times D$  where  $N$ = batch size,  $H$ =image height,  $W$ =image width and  $D$ = image depth (number of channels)
- $y_{i,h,w,d}$  is the corresponding element of the output tensor
- $\gamma$  and  $\beta$  are trainable parameters called *scale* and *center*
- $\epsilon$  is a small constant to prevent numerical instability (i.e. divide by zero)

# BN Practise

- Full BN is performed only in training
- Global standard deviations and means of the training set  $\sigma_d^g$  and  $\mu_d^g$  are stored during training (as the moving average of  $\sigma_d$  and  $\mu_d$  of each batch)
- In testing (inference) the stored global standard deviations and means are used.

# Other Issues

- How to select hyper-parameters
  - Learning rate
  - Regularization parameters
  - Model architecture
- Initialization of model parameters