

Introduction to Tensorflow

Narada Warakagoda

Modes

Tensors

Model

Loss Opti Data-Feed

Training

1 / 38

What is Tensorflow?

- An open source software library from Google.
- Supports deep learning and some machine learning algorithms.
- Can be run on multiple machines and GPUs.
- Several APIs Python, Java and C++.
- Tensorflow Python API is tightly integrated with Keras, a high level set of APIs for deep learning.
- Latest version (as of July 2022) is 2.9
- For more information refer https://www.tensorflow.org/api_docs/python/tf

Modes

Tensors

Model

Loss Opti

Data-Feed

Training

2 / 38

Tensorflow execution modes


- Graph execution mode
 - Consists of two phases:
 - Creation of a computation graph
 - Run the data through the computation graph
 - Natural way of supporting Automatic Differentiation
 - Early Tensorflow versions supported this mode only
- Eager execution mode
 - No execution phases similar to other languages such as plain Python
 - Useful in debugging
 - Generally slower than Graph execution mode
 - Default mode in Tensorflow 2.x

Graph mode example

- Consider a very simple model $\hat{\mathbf{y}} = f(\mathbf{x}\mathbf{W} + \mathbf{b})$
- Loss function $C = g(\hat{\mathbf{y}}, \mathbf{y})$

where

- $\mathbf{x} = [x_1, x_2, \dots, x_L]$ is the input training example
- $\mathbf{y} = [y_1, y_2, \dots, y_K]$ is the label of the example
- \mathbf{W} is the weight matrix of the network
- \mathbf{b} is the bias vector of the network
- $f(\cdot)$ is the activation function and assume that it is a ReLU
- $g(\cdot)$ is the loss function, and assume that it represents cross entropy Xent



Modes

Tensors

Model


Loss Opti

Data-Feed

Training

Optimization with the computation graph

- Gradient computation graph is automatically generated (i.e. Automatic Differentiation)



User defined
forward computation graph

Tensorflow defined
gradient computation graph

Advantages of Graph mode

- Flexibility (i.e. not dependent on Python)
- Can be exploited on different devices
- Used in exporting and importing models
- Can easily be optimized

Modes

Tensors

Model

Loss Opti

Data-Feed

Training

6 / 38

Enforcing Graph mode

- Use `tf.function` directly or as a Python decorator

```
1 # Define a Python function.
2 def a_regular_function(x, y, b):
3     x = tf.matmul(x, y)
4     x = x + b
5     return x
6
7 # 'a_function_that_uses_a_graph' is a TensorFlow 'Function'.
8 a_function_that_uses_a_graph = tf.function(a_regular_function)
```

```
1 # Define a Python function. with decoration
2 @tf.function
3 def a_regular_function(x, y, b):
4     x = tf.matmul(x, y)
5     x = x + b
6     return x
7
8 # 'a_regular_function' is a TensorFlow 'Function'.
```

Eager Execution mode

- This is the default mode of Tensorflow
- Code is executed line by line
- Since there is no graph, Tensorflow needs to “record” all relevant operations to support Automatic Differentiation
 - `tf.GradientTape` is the data structure used for this purpose.

Tensorflow Tensors


- Main data structure of Tensorflow is a Tensor
- Tensors can be created using different functions
 - `tf.constant`
 - `tf.variable`
 - `tf.zeros`
 - `tf.ones` and more ..

A scalar, shape: []

4


`tf.constant(4)`

A vector, shape: [3]



`tf.constant([2.0,3.0,4.0])`

A matrix, shape: [3, 2]



`tf.constant([[1,2],
[3,4],
[5,6]])`

Modes

Tensors

Model


Loss Opti

Data-Feed

Training

Tensorflow rank-3 Tensor

A 3-axis tensor, shape: [3, 2, 5]



```
rank_3_tensor = tf.constant([  
    [[0, 1, 2, 3, 4],  
     [5, 6, 7, 8, 9]],  
    [[10, 11, 12, 13, 14],  
     [15, 16, 17, 18, 19]],  
    [[20, 21, 22, 23, 24],  
     [25, 26, 27, 28, 29]],  
])
```

Modes

Tensors

Model

Loss Opti


Data-Feed

Training


10 / 38

Tensorflow rank-4 Tensor

A rank-4 tensor, shape: [3, 2, 4, 5]



```
tf.zeros([3, 2, 4, 5])
```



Main Steps in training a model with Tensorflow

- Define the model architecture
- Define the loss function
- Choose an optimizer and learning rate
- Create the data feeding code
- Write the training/validation loop
 - Training step
 - Run the next training mini-batch through the model
 - Calculate the loss
 - Calculate the gradients and update model parameters
 - Update loss (and metric) summaries
 - Validation step
 - Run the next validation mini-batch through the model
 - Calculate the loss
 - Update loss (and metric) summaries

Modes

Tensors

Model

Loss Opti Data-Feed

Training

12 / 38



Model Architecture (Layers)

- Typically models are composed using `tf.keras.layers` which can be
 - Predefined layers (eg: `tf.keras.layers.Conv2D` , `tf.keras.layers.Dense`)
 - Custom layers (create your own layers by extending `tf.keras.layers.Layer`)

Modes

Tensors

Model

Loss Opti

Data-Feed

Training

13 / 38

Pre-defined Layers Example (2D Convolution)

```
1 conv = tf.keras.layers.Conv2D(  
2     filters, # e.g. 64  
3     kernel_size, # e.g. (3, 3)  
4     strides=(1, 1),  
5     padding='valid', # other option is 'same'  
6     data_format=None,  
7     dilation_rate=(1, 1),  
8     activation=None,  
9     use_bias=True,  
10    kernel_initializer='glorot_uniform',  
11    bias_initializer='zeros',  
12    kernel_regularizer=None,  
13    bias_regularizer=None,  
14    activity_regularizer=None,  
15    kernel_constraint=None,  
16    bias_constraint=None,  
17    **kwargs)
```

Predefined layer example (Dropout)

```
1 dropout = tf.keras.layers.Dropout(0.5)
2 tf.random.set_seed(123)
3 x = [[1.0, 1.0, 1.0], [1.0, 1.0, 1.0]]
4 dropout(x, training=True) # [[0., 2., 0.], [2., 2., 2.]]
5 dropout(x, training=True) # [[2., 0., 2.], [2., 2., 0.]]
6 dropout(x, training=False) # [[1., 1., 1.], [1., 1., 1.]]
```

Pre-defined layer example (Batch Normalization)

```
1 batch_norm = tf.keras.layers.BatchNormalization(axis=-1,  
  ↪ center=False, scale=False)  
2 x = np.array([[ -1.0, 4.0, 1.0], [1.0, -4.0, 3.0]])  
3 # [batch_size, num_features] == [2, 3]  
4 # feature 1: mean = 0, std ~ = 1  
5 # feature 2: mean = 0, std ~ = 4  
6 # feature 3: mean = 2, std ~ = 1  
7 batch_norm(x, training=True) # ~ = [[-1.0, 1.0, -1.0], [1.0,  
  ↪ -1.0, 1.0]]  
8 batch_norm(x, training=False) # ~ = ?
```

Keras custom layer (example)

```
1 class FullyConnected_v2(tf.keras.layers.Layer):
2     def __init__(self, num_outputs):
3         super(FullyConnected_v2, self).__init__()
4         self.num_outputs = num_outputs
5
6     def build(self, input_shape):
7         """Assume input_shape[0] is batch size and input_shape[1] is size
8         ↪ of input samples."""
9         self.W = tf.Variable(np.random.uniform(-0.1, 0.1,
10         ↪ size=[input_shape[1], self.num_outputs]))
11         self.b = tf.Variable(np.zeros(self.num_outputs))
12
13     def call(self, x):
14         return tf.matmul(x, self.W) + self.b
15
16 # array of shape [batch_size, num_inputs] ==> [batch_size, 5]
17 fc = FullyConnected_v2(num_outputs=5)
18 # array of shape [2, 3] ==> [2, 5]
19 fc(np.array([[1.0, 0.4, 0.2], [-0.4, 0.3, 0.2]]))
```

Notes on custom layers

- The build method is called the first time `__call__` is called.
- We implement `call` rather than `__call__`. The `__call__` method is implemented in the parent class, and will call 'call'
- We may also use `self.add_variable` method to add variables to layer
- Remember to call super method in `__init__` to initialize Layer class properly.
- Optional: Implement `get_config` and `compute_output_shape` (for serialization and model summary purposes respectively).

Model Architecture (Models)

- Once you have all the layers, compose the intended model using one of
 - Keras sequential API
 - Keras Functional API
 - Extending `tf.keras.Model`

Modes

Tensors

Model

Loss Opti

Data-Feed

Training

19 / 38

Composing model (sequential API)

- Just add your layers to the `tf.keras.models.sequential` call

```
1 from tensorflow.keras.layers import Conv2D, Flatten, Dense
2
3 model = tf.keras.models.Sequential([
4     Conv2D(32, kernel_size=3, activation='relu'),
5     Conv2D(64, kernel_size=3, strides=2, activation='relu'),
6     Flatten(),
7     Dense(128, activation='relu'),
8     Dense(10, activation='softmax')
9 ])
```

Composing model (Functional API)

- Need to define a dummy input (specify the size without batch size)
- Then call the layers one after other

```
1 inputs = tf.keras.Input(shape=(32, 32, 3))
2 x = Conv2D(32, kernel_size=3, activation='relu')(inputs)
3 x = Conv2D(64, kernel_size=3, strides=2, activation='relu')(x)
4 x = Flatten()(x)
5 x = Dense(128, activation='relu')(x)
6 outputs = Dense(10, activation='softmax')(x)
7
8 model = tf.keras.Model(inputs=inputs, outputs=outputs)
```

Composing model (Extending `tf.keras.Model`)

- Need to implement `__init__()` and `call()` methods.

```
1 class MyModel(tf.keras.Model):  
2     def __init__(self):  
3         super(MyModel, self).__init__()  
4         self.conv1 = Conv2D(32, kernel_size=3, activation='relu')  
5         self.conv2 = Conv2D(64, kernel_size=3, strides=2, activation='relu')  
6         self.flatten = Flatten()  
7         self.d1 = Dense(128, activation='relu')  
8         self.d2 = Dense(10, activation='softmax')  
9  
10    def call(self, x):  
11        x = self.conv1(x)  
12        x = self.conv2(x)  
13        x = self.flatten(x)  
14        x = self.d1(x)  
15        return self.d2(x)  
16 model = MyModel()
```

Loss function

- `tf.keras.losses` has many pre-defined loss functions to choose from
 - `MeanSquaredError` (MSE)
 - `SparseCategoricalCrossEntropy` (CE when labels are sparse)
 - `CategoricalCrossEntropy` (CE when labels are one-hot encoded)

```
1 mse = tf.keras.losses.MeanSquaredError()
2 loss = mse([0., 0., 1., 1.], [1., 1., 1., 0.])
3 print('Loss: ', loss.numpy())    # Loss: 0.75
4
5 cce = tf.keras.losses.SparseCategoricalCrossEntropy(from_logits=
    ↪ =False) # set true if not softmax
    ↪ applied
6 loss = cce([0, 1, 2], [[.9, .05, .05], [.5, .89, .6], [.05,
    ↪ .01, .94]])
7 print('Loss: ', loss.numpy())    # Loss: 0.3239
```

Optimizer

- tf.keras.optimizers has several optimizers to choose from. Pick one of them and create an optimizer object.
- Optimizer implements an update rule only (Gradient computation is outside of its scope)

```
1 import tensorflow as tf
2
3 # Stochastic gradient descent, with or without momentum
4 op1=tf.keras.optimizers.SGD(learning_rate=0.01, momentum=0.0,
5                               nesterov=False, name='SGD', **kwargs)
6
7 op2=tf.keras.optimizers.RMSprop(learning_rate=0.001, rho=0.9,
8                                   momentum=0.0, epsilon=1e-07,
9                                   centered=False, name='RMSprop', **kwargs)
10
11 op3=tf.keras.optimizers.Adam(learning_rate=0.001, beta_1=0.9,
12                               beta_2=0.999, epsilon=1e-07,
13                               amsgrad=False, name='Adam', **kwargs)
```

Data Feeding

- Typical data feeding pipeline is based on `tf.data.Dataset`
- Basic steps:
 - Create a `tf.data.Dataset` object with the data you have. It can “consume” several different sources such as
 - numpy arrays,
 - Python generators,
 - tfrecord data,
 - CSV data and
 - sets of files
 - Perform desired operations on the `Dataset` object (eg: `shuffle`, `map`, `repeat` and `batch`)
 - Iterate through the `Dataset` object to get data for training (or validation).

Modes

Tensors

Model

Loss Opti

Data-Feed

Training

25 / 38

Data feeding example

- Dataset object consumes a Python generator
- Functions used in code are in the next slide

```
1
2 #####
3 # Create a Dataset object consuming Python generator
4 dataset = tf.data.Dataset.from_generator(
5     generator_for_filenames(image_names, class_labels),
6     output_types=(tf.string, tf.string),
7     output_shapes=(None, None)
8 )
9
10 # Shuffle data
11 dataset = dataset.shuffle(buffer_size=100)
12
13 # Map Dataset objects (file names to images)
14 dataset = dataset.map(read_image)
15
16 #Map Dataset objects (images to preprocessed images)
17 dataset = dataset.map(preprocess)
18
19 #Introduce batching
20 dataset = dataset.batch(batch_size)
21
22
23
24
25 #####
26 # Now you can use the Dataset
27 for image, label in dataset:
28     print(image, label)
29     #Feed your data through the model etc.
30     # pred=model(image)
```

Modes

Tensors

Model

Loss Opti

Data-Feed

Training

26 / 38

Data feeding example (Auxiliary functions)

```
1
2
3 def generator_for_filenames(*filenames):
4     #Wrapping a list of filenames as a generator function
5     def generator():
6         for f in zip(*filenames):
7             yield f
8     return generator
9
10
11 def preprocess(image, class_lab):
12     #A preprocess function the is run after images are read.
13     #Here you can do augmentation and other pre-processing
14     #on the images.
15
16     # Set images size to a constant
17     image = tf.image.resize(image, [HEIGHT, WIDTH])
18
19     # Other pre-processing stuff here
20     return image, class_lab
21
22
23 def read_image(img_f, class_lab):
24     #Read images from file using tensorflow
25     img_reader = tf.io.read_file(img_f)
26     img = tf.image.decode_png(img_reader, channels=3)
27
28     return img, class_lab
```

Modes

Tensors

Model

Loss Opti **Data-Feed**

Training

27 / 38

Training Loop

- There are two main alternatives
 - Use the predefined methods `tf.keras.model.compile` and `tf.keras.model.fit`
 - Write training steps using non-Keras (Pure Tensorflow) API
 - Run the next training mini-batch through the model
 - Calculate the loss
 - Calculate the gradients and update model parameters
 - Update loss (and metric) summaries
- Previous steps are common to any of these alternatives (i.e. definition of model, loss, optimizer and data feeder etc.)

Modes

Tensors

Model

Loss Opti Data-Feed

Training

28 / 38



Common steps

```
1 import tensorflow as tf
2 from tensorflow.keras import losses, metrics, optimizers
3
4
5 #define your model in get_model() function
6 model = get_model()
7
8 #define your loss function
9 loss_fn = losses.Crossentropy(from_logits=False)
10
11 #define your optimizer
12 optimizer_fn = optimizers.Adam(lr=1e-4)
13
14 #Create a Dataset object in get_train_data() function
15 train_dataset = get_train_data()
16
17 #create a Dataset object in get_val_data() function
18 val_dataset = get_val_data()
19
20 # Create function to calculate an optional custom metric
21 metric_fn = get_metric_fn()
```

Keras compile-fit

```
1 model.compile(  
2     optimizer=optimizer_fn, # Optimizer  
3     loss=loss_fn,# Loss function to minimize  
4     metrics=[metric_fn],# List of metrics to monitor  
5 )  
6  
7 history = model.fit(  
8     train_dataset,  
9     batch_size=64,  
10    epochs=2,  
11    validation_data=val_dataset,  
12 )
```

Modes

Tensors

Model

Loss Opti Data-Feed

Training

30 / 38



Pure Tensorflow training loop

```
1 # Start train/validation loop, assume train_epochs=100
2 for epoch in range(train_epochs):
3
4     for image, y in train_dataset:
5         y_pred = train_step(image, y)
6         step += 1
7
8
9     # reset metrics
10    train_loss.reset_states()
11    train_accuracy.reset_states()
12
13    # Do validation after each epoch
14    for image, y in val_dataset:
15        y_pred = val_step(image, y)
16
17    # reset metrics
18    val_loss.reset_states()
19    val_accuracy.reset_states()
```

Train Step

Validation Step

Modes

Tensors

Model

Loss Opti Data-Feed

Training

31 / 38

Pure Tensorflow training step

```
1  # Used to running averages for summaries
2  train_accuracy = metrics.Mean()
3  train_loss = metrics.Mean()
4  val_accuracy = metrics.Mean()
5  val_loss = metrics.Mean()
6
7
8  # Train step for a single batch
9  def train_step(image, y):
10
11      # get predictions loss and update model parameters
12      with tf.GradientTape() as tape:
13          y_pred = model(image, training=True)
14          loss = loss_fn(y, y_pred)
15          grads = tape.gradient(loss, model.trainable_variables)
16          optimizer_fn.apply_gradients(zip(grads, model.
17                                     trainable_variables))
18
19      # Update running averages for summaries
20      m = metric_fn(y, y_pred)
21      train_loss.update_state(loss)
22      train_accuracy.update_state(m)
23
24      return y_pred
```

Modes

Tensors

Model

Loss Opti Data-Feed

Training

32 / 38



Pure Tensorflow validation step

```
1 #validation step for a single batch
2 def val_step(image, y):
3
4     y_pred = model(image)
5     loss = loss_fn(y, y_pred)
6
7     # Update running averages for summaries
8     m = metric_fn(y, y_pred)
9     val_loss.update_state(loss)
10    val_accuracy.update_state(m)
11
12    return y_pred
```

Modes

Tensors

Model

Loss Opti Data-Feed

Training

33 / 38



Tensorflow RNN implementation

- Keras built-in RNN networks (standard RNN configurations)
 - `tf.keras.layers.SimpleRNN`
 - `tf.keras.layers.LSTM`
 - `tf.keras.layers.GRU`
- RNN Cell wrapped in `tf.keras.layers.RNN`
 - `tf.keras.layers.SimpleRNNCell`
 - `tf.keras.layers.LSTMCell`
 - `tf.keras.layers.GRUCell`
 - Custom cell extending `tf.keras.layers.AbstractRNNCell`

background

RNN Cells

Configs

LSTM

Variants


–

Implement

34 / 38

Keras built-in RNNs (I)


```
1 import tensorflow as tf
2 import numpy as np
3
4 inputs = np.random.random([32, 10, 8]).astype(np.float32)
5
6 # RNN with only last output
7 simple_rnn1 = tf.keras.layers.SimpleRNN(4)
8
9 # The output has shape '[32, 4]'.
10 output = simple_rnn1(inputs)
11
12 # RNN with output at all time steps
13 simple_rnn2 = tf.keras.layers.SimpleRNN(4, return_sequences=True,
14                                     return_state=True)
15
16 # whole_sequence_output has shape '[32, 10, 4]'.
17 # final_state has shape '[32, 4]'.
18 whole_sequence_output, final_state = simple_rnn2(inputs)
19
20 # Bidirectional RNN based on simple_rnn2
21 simple_bidirectional = tf.keras.layers.Bidirectional(simple_rnn2)
```



Keras built-in RNNs (II)

- Encoder-decoder using LSTMs

```
1 import tensorflow as tf
2 import numpy as np
3
4 # Define encoder and decoder values for running the model
5 enc_input_val = np.random.random([32, 10, 8]).astype(np.float32)
6 dec_input_val = np.random.random([32, 6, 7]).astype(np.float32)
7
8 #Define the encoder
9 encoder_input = tf.keras.layers.Input(shape=(10,8))
10 encoder=tf.keras.layers.LSTM(64, return_state=True, name="encoder")
11
12 # Return states in addition to output
13 output, state_s, state_c = encoder(encoder_input)
14 encoder_state = [state_s, state_c]
15
16 # Evaluate encoder
17 output_val, state_s_val, state_c_val = encoder(encoder_input_val)
18 print(output_val.shape, state_s_val.shape, state_c_val.shape)
19
20 #define the decoder
21 decoder_input = tf.keras.layers.Input(shape=(6,7))
22 decoder = tf.keras.layers.LSTM(64, return_sequences=True,
23                               name="decoder")
24
25 # Pass the 2 states to a new LSTM layer, as initial state
26 dec_output = decoder(decoder_input, initial_state=encoder_state)
27 output = tf.keras.layers.Dense(10)(dec_output)
28
29 #define the full model (encoder-decoder)
30 model = tf.keras.Model([encoder_input, decoder_input], output)
31 #model.summary()
32
33 # Predict using the model
34 output_val= model.predict([encoder_input_val, decoder_input_val])
35 print(output_val.shape)
```



background

RNN Cells

Configs

LSTM


Variants

Implement



36 / 38

Cells wrapped in keras.RNN

```
1 inputs = np.random.random([32, 10, 8]).astype(np.float32)
2
3 # RNN wrapping a simple RNN cell
4 rnn = tf.keras.layers.RNN(
5     tf.keras.layers.SimpleRNNCell(4),
6     return_sequences=True,
7     return_state=True)
8
9 # whole_sequence_output has shape '[32, 10, 4]'.
10 # final_state has shape '[32, 4]'.
11 whole_sequence_output, final_state = rnn(inputs)
```



```
1 inputs = tf.random.normal([32, 10, 8])
2 # RNN wrapping an LSTM cell, single output
3 rnn_single = tf.keras.layers.RNN(tf.keras.layers.LSTMCell(4))
4 output = rnn_single(inputs)
5 print(output.shape)
6
7 # RNN wrapping an LSTM cell, sequential output
8 rnn_seq = tf.keras.layers.RNN(
9     tf.keras.layers.LSTMCell(4),
10    return_sequences=True,
11    return_state=True)
12 whole_seq_output, final_state_s, final_state_c = rnn_seq(inputs)
```



Custom Cells wrapped in keras.RNN

- Class extends `tf.keras.layers.AbstractCell`
- Call method has the signature `(output, next_state) = call(input, state)`

```
1
2 class MyVanillaRNNCell(tf.keras.layers.AbstractRNNCell):
3     def __init__(self, units):
4         super(MyVanillaRNNCell, self).__init__()
5         self.units = units
6         self.dense = layers.Dense(units)
7
8     @property
9     def state_size(self):
10         return self.units
11
12
13     def call(self, x, state):
14         c = tf.concat([x, state], axis=-1)
15         h = self.dense(c)
16         output = activations.tanh(h)
17         return output, output
18
19
20 my_own_rnn=tf.keras.layers.RNN(MyVanillaRNNCell (4))
```

